

# *Data representation*

Floating-point numbers and multiple-precision

# Data representation

## Real numbers

Example of exact representations:

$$123.456 = 123456 \times 10^{-3} = 1.23456 \times 10^2$$

significand x base<sup>exponent</sup>

## Integers

42 : what is the meaning of all this?

how does the computer store these numbers?

what are the limitations?


# 8-bits, 16-bits, 32-bits, 64-bits, 128-bits ... ?

(nes, snes, ps, n64, ps2, xbox, ... ;-)

Each CPU implements native **n-bits operations** [ current PCs: 64-bits ]

Example: it limits the memory allocation and data representation

A 3-bit memory controller  
can only **point** to  $2^3 = 8$  addresses



A 3-bit memory bank  
can only **represent** 8 abstract entities


000 = apple  
001 = avocado  
010 = banana  
011 = lemon  
100 = pineapple  
101 = strawberry  
110 = kiwi  
111 = grape



# Integer representation

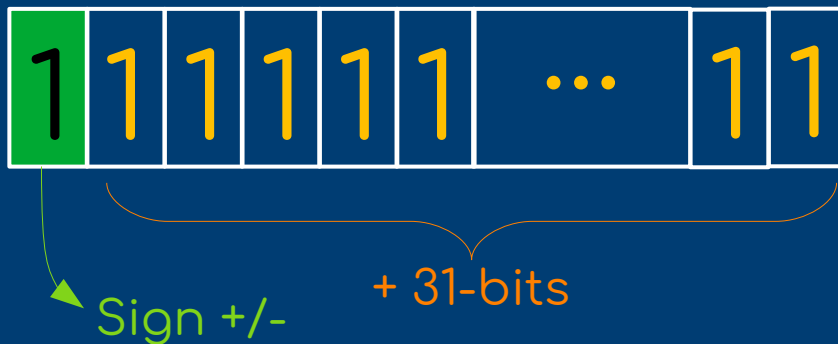
From the previous slide, a 3-bit computer can only represent 8 integers

3-bit binary memory


$$1 \times 2^2 + 0 \times 2^1 + 1 \times 2^0 = 5$$

Min: (000)<sub>b</sub> = 0  
Max: (111)<sub>b</sub> = 7

Signed 32-bit integer standard



$$= (-1)^{b_{31}} \times (b_{30} \times 2^{30} + b_{29} \times 2^{29} + \dots + b_0 \times 2^0)$$

“sign-magnitude representation”

Range 9+ digits

$$\begin{aligned} \text{Min: } & -2^{31} = -2,147,483,648 \\ \text{Max: } & +2^{31}-1 = +2,147,483,647 \end{aligned}$$

“2’s complement representation”

# Real numbers : floating point

significand x base<sup>exponent</sup>

Signed 32-bit float standard



While the set of real numbers is dense / continuous...  
only a limited number of entities can be represented on a finite structure  
→ loss of precision

# Comparing real numbers

Let us try this in C

1) initialize the variables as

```
x = 0.1  
y = 3*x  
z = 0.3
```

2) check the output of the equality tests

```
if (z == 0.3) { ... }  
if (z == y) { ... }  
If (z == 3*0.1) { ... }
```

3) What is happening? Let us print x, y, z with many digits

```
printf('x = 0.30f\n', x)  
printf('y = 0.30f\n', y)  
printf('z = 0.30f\n', z)
```

# Why is $0.1 = 0.100000000000000000005551115123126$ ?

If a float is simply stored as **significand** x **base** <sup>exponent</sup>

shouldn't we have  $0.1 = 1 \times 10^{-1}$  ?

... but computers use **base 2** !!!

## Examples

→ the fraction  $1/10$  in base 10 is 0.1 (exact)

→ the fraction  $1/3$  in base 10 is  $0.\overline{3} = 0.333333... \sim 0.333334$

→ the fraction  $1/10$  in base 2 is  $0.0\overline{0011} = 0.00011001100110011...$

An exact fraction in base 10 might be a repeating fraction in base 2

Truncation yields:  $(0.0001100110011)_2 = (0.0999755859375)_{10}$

# Multiple, or arbitrary-precision

Default types:

C

---

[int] = 32 bits  
[float] = 32 bits  
[double] = 64 bits  
...

Python

---

[int] = arbitrary  
[float] = 64 bits

Allows you to freely choose the precision

Floats: 32 bits ~ 7 digits, 64 bits ~ 16 digits, 128 bits ~ 34 digits

---

GMP: GNU Multiple Precision Arithmetic Library

[<https://gmplib.org/>]

MPFR: multiple-precision floating-point computations  
with correct rounding

[<https://www.mpfr.org/>]

---

Python mpmath [<http://mpmath.org/>]

→ for multiple-precision in Python, uses GMP if available

Python SymPy [<https://www.sympy.org/>]

→ for **Symbolic** mathematics in Python



# Example in Python

From “Why and How to Use Arbitrary Precision”,  
Ghazi et al, Computing in Science & Engineering 12, 62-65 (2010)

Let us try to calculate  $d = 173746a + 94228b - 78487c$ , with

$$a = \sin(10^{22}), b = \log_{10}(17.1), c = \exp(0.42)$$

Standard Python

```
import numpy as np
a = np.sin(1e22)
b = np.log(17.1)
c = np.exp(0.42)
d = 173746*a + 94228*b - 78487*c
print('d =', d)
      2.9103830456733704e-11
```

Exact value:  
 $-1.341818958... \times 10^{-12}$

## The same Python example, but using **mpmath**

```
from mpmath import mp
mp.dps = 40 # defines the decimal precision
print(mp) # to check again
```

```
a = mp.sin('1e22')
b = mp.log('17.1')
c = mp.exp('0.42')
d = 173746*a + 94228*b - 78487*c
print('d = ', d)
      -1.34181895782961046706215258814e-12
```

Exact value:  
 $-1.341818958... \times 10^{-12}$

---

Notice that the **real numbers are set by strings**, otherwise Python would first convert to floating-point and lose precision. TRY IT!

# Overall features of these libraries

	python numpy+scipy	python mpmath	C+GMP+MPFR
Float precision	64 bits	arbitrary	arbitrary
Trigonometric functions + Special functions	<code>np.sin(...)</code> <code>np.exp(...)</code> <code>sp.special.jv(...)</code> <code>sp.special.zeta(...)</code>	<code>mp.sin(...)</code> <code>mp.exp(...)</code> <code>mp.besselj(...)</code> <code>mp.zeta(...)</code>	<code>mpfr_sin(...)</code> <code>mpfr_exp(...)</code> <code>mpfr_jn(...)</code> <code>mpfr_zeta(...)</code>
Numerical calculus	Root finding, sum, quadrature (integrals), differentiation, ODE (RK4), Taylor, Fourier, ...		Check other libraries:  GSL  mpack = blas+lapack
Linear algebra	Matrix/vector operations (products, inverse, determinant), SVD, linear systems, eigenproblems, matrix functions (exp, cos, ...), ...		Boost C++ <a href="http://www.boost.org">[www.boost.org]</a>

# Exercise: factorial

Let us go back to our factorial implemented on a for loop:

```
def myfactorial(n):  
    f = 1  
    for i in range(1, n+1):  
        f *= i  
    return f
```

**CENSORED**

- 1) Try to run it for  $n=30$   
→ expected result: 265252859812191058636308480000000
- 2) Now, change your implementation to have the initial  $f$  as a float:  
→  $f = 1.0$
- 3) And check again for  $n=10$  and  $n=30$ . Float → **loss of precision!**
- 4) Change the code to initialize  $f$  as a multiple precision float